Table of contents

Introduction **4**
Primary Data Types and OS Allocation **4**
i) Data Type Sizes Across Architectures 4
ii) Practical Example: Observing Data Type Sizes 5
OS Memory Segments **5**
i) Memory Segment Overview 5
ii) Memory Layout Visualization 6
Practical Code Examples **6**
i) Observing variable addresses 6
ii) Pointer and memory allocation 7
Laboratory Exercises **7**
i) Lab 1: Data Type Size Exploration 7
ii) Lab 2: Memory Segment Address Mapping 8
Summary and Key Takeaways **11**
Knowledge Check Question with Answers **12**

List of Figures

Figure 1: Program to observe data type sizes written inside file1.c..... 5
Figure 2: Process memory layout (high to low addresses).....6
Figure 3: Program to observe variable addresses written inside file2.c.....6
Figure 4: Correct pointer and heap usage..... 7
Figure 5: Linux command 'touch file1.c' to create file1.c file and 'nano file1.c' to write C program inside file1.c.....7
Figure 6: Linux command 'cat file1.c' to view the contents inside file1.c..... 7
Figure 7: Linux command of 4 compilation stages to convert file1.c to executable file "file1"..... 8
Figure 8: Linux command './file1' to show the program output and 'echo $?' to show successful execution of the program.....8
Figure 9: Linux command 'touch file2.c' to create file2.c and 'nano file2.c' to write C program on that file2.c.....8
Figure 10: Linux command 'cat file2.c' to view the contents.....9
Figure 11: Linux command of 4 compilation stages to convert file2.c to executable file "file2"..... 9
Figure 12: Linux command './file2' to view the program output and 'echo $?' to view successful execution of the program.....9
Figure 13: Linux command 'touch file3.c' to create file3.c and 'nano file3.c' to write contents inside file3.c.....10
Figure 14: Linux command 'cat file3.c' to view the contents inside file3.c..... 10
Figure 15: Linux command of 4 compilation stages to convert file3.c to executable file "file3"..... 10
Figure 16: Linux command './file3' to show the program output and 'echo $?' to view successful execution of the program and 'ls' to view a list of files.....10

List of Table

Table I: Datatype sizes across system architecture..... 4
Table II: Process memory segments and their purposes..... 5

Introduction

Data types form the foundation of programming and operating system interaction. A data type is a contract between programmer, compiler, and operating system that specifies:
1) The number of bytes to allocate in memory.
2) How to interpret those bytes (integer, character, floating point).
3) What operations are valid on that data.
When a programmer declares a variable, the OS must allocate physical or virtual memory to store that data. Understanding this relationship is crucial for writing efficient, secure, and portable code.

Primary Data Types and OS Allocation

i) Data Type Sizes Across Architectures

Table I shows how data type sizes vary across different system architectures, directly impacting OS memory allocation.

Data Type	16-bit	32-bit	64-bit	OS Impact
char	1 byte	1 byte	1 byte	Consistent across architectures
short	2 bytes	2 bytes	2 bytes	Memory-efficient counters
int	2 bytes	4 bytes	4 bytes	Affects pointer arithmetic
long	4 bytes	4 bytes	8 bytes	Critical for 64-bit addressing
long long	8 bytes	8 bytes	8 bytes	Large file offsets
float	4 bytes	4 bytes	4 bytes	IEEE 754 standard
double	8 bytes	8 bytes	8 bytes	Precision calculations
pointer	2 bytes	4 bytes	8 bytes	Address space limit

Table I: Datatype sizes across system architecture

ii) Practical Example: Observing Data Type Sizes

```
  GNU nano 8.7.1                  file1.c *
#include <stdio.h>

int main() {
        printf("=== Data Type Sizes on This System ===\n\n");

        printf("char: %zu bytes\n", sizeof(char));
        printf("int: %zu bytes\n", sizeof(int));
        printf("float: %zu bytes\n", sizeof(float));
        printf("double: %zu bytes\n", sizeof(double));
        printf("long: %zu bytes\n", sizeof(long));
        printf("pointer: %zu bytes\n", sizeof(int*));

        return 0;
}
```

Figure 1: Program to observe data type sizes written inside file1.c.

OS Memory Segments

The operating system creates a process address space, a virtual view of memory divided into segments. Each segment serves a specific purpose and has different characteristics.

i) Memory Segment Overview

Segment	Purpose	Grows	Managed By
Text	Machine code (instructions)	Fixed	Compiler/Loader
Data	Initialized global/static variables	Fixed	Compiler/OS
BSS	Uninitialized global/static variables	Fixed	OS
Heap	Dynamically allocated memory	Upward	Programmer + OS
Stack	Local variables, function calls	Downward	Compiler + OS

Table II: Process memory segments and their purposes.

ii) Memory Layout Visualization

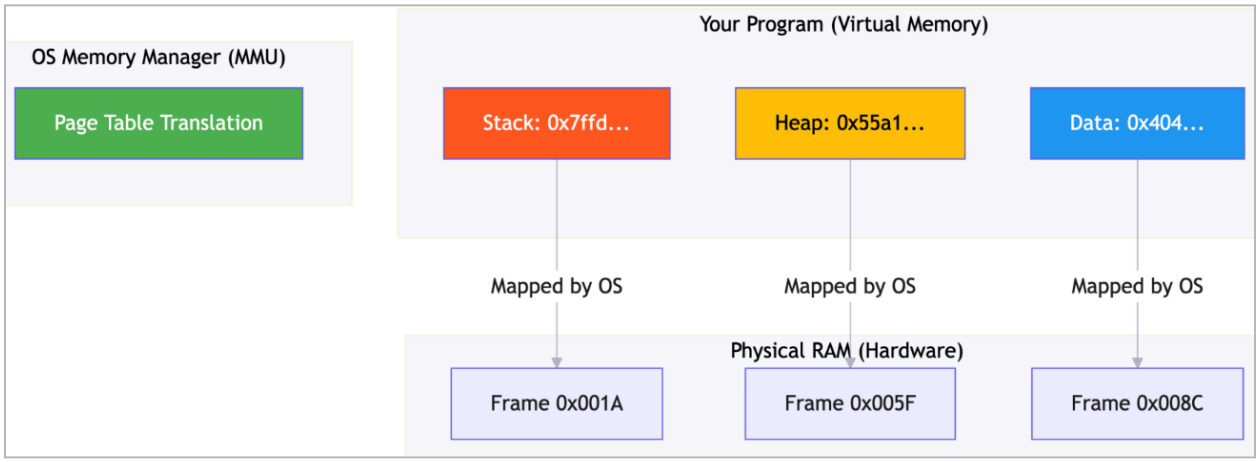


Figure 2: Process memory layout (high to low addresses).

Key Observation: Stack addresses are high (0x7ffd...), Heap addresses are mid-range (0x55a1...), and BSS/Data/Text addresses are low (0x404...).

Practical Code Examples

i) Observing variable addresses

```
  GNU nano 8.7.1                                  file2.c *
#include <stdio.h>
#include <stdlib.h>

int global_var = 150; // Data segment
int bss_var; // BSS/Data segment

int main() {
        int local_var = 30; // Stack
        static int local_static = 20; // Data segment

        int *heap_var = (int *)malloc(sizeof(int)); //Heap*heap_var = 500;
        printf(".......Variable Addresses.......\n\n");
        printf("global_var: %p (Data/BSS)\n", (void*)&global_var);
        printf("bss_var: %p (Data/BSS)\n", (void*)&bss_var);
        printf("local_var: %p (Stack)\n", (void*)&local_var);
        printf("local_static: %p (Data/BSS)\n", (void*)&local_static);
        printf("heap_var ptr: %p (Stack - pointer)\n",(void*)&heap_var);
        printf("*heap_var: %p (Heap - data)\n", (void*)heap_var);

        free(heap_var);
        return 0;
}
```

Figure 3: Program to observe variable addresses written inside file2.c.

ii) Pointer and memory allocation

```
 GNU nano 8.7.1                                                  file3.c
#include <stdio.h>
#include <stdlib.h>

int main() {
        int *forheap; // Declare pointer (lives on STACK)
        forheap = (int *)malloc(sizeof(int)); //Allocate memory on HEAP first!
        *forheap = 30; // store value to allocated memory// Print addresses
        printf("Address of pointer (STACK): %p\n", (void*)&forheap);
        printf("Address of data (HEAP): %p\n", (void*)forheap);
        printf("Value stored: %d\n", *forheap);

        return 0;
}
```

Figure 4: Correct pointer and heap usage.

Laboratory Exercises

i) Lab 1: Data Type Size Exploration

```
root@DESKTOP-OAP8OK0:/batch38/Week4# touch file1.c
root@DESKTOP-OAP8OK0:/batch38/Week4# nano file1.c
```

Figure 5: Linux command 'touch file1.c' to create file1.c and 'nano file1.c' to write C program inside file1.c

```
root@DESKTOP-OAP8OK0:/batch38/Week4# cat file1.c
#include <stdio.h>

int main() {
        printf("=== Data Type Sizes on This System ===\n\n");

        printf("char: %zu bytes\n", sizeof(char));
        printf("int: %zu bytes\n", sizeof(int));
        printf("float: %zu bytes\n", sizeof(float));
        printf("double: %zu bytes\n", sizeof(double));
        printf("long: %zu bytes\n", sizeof(long));
        printf("pointer: %zu bytes\n", sizeof(int*));

        return 0;
}
```

Figure 6: Linux command 'cat file1.c' to view the contents inside file1.c.

```
root@DESKTOP-OAP8OK0:/batch38/Week4# gcc -E file1.c -o file1.i
root@DESKTOP-OAP8OK0:/batch38/Week4# gcc -S file1.i -o file1.s
root@DESKTOP-OAP8OK0:/batch38/Week4# gcc -c file1.s -o file1.o
root@DESKTOP-OAP8OK0:/batch38/Week4# gcc file1.o -o file1
```

Figure 7: Linux command of 4 compilation stages to convert file1.c to executable file "file1".

```
root@DESKTOP-OAP8OK0:/batch38/Week4# ./file1
=== Data Type Sizes on This System ===

char: 1 bytes
int: 4 bytes
float: 4 bytes
double: 8 bytes
long: 8 bytes
pointer: 8 bytes
root@DESKTOP-OAP8OK0:/batch38/Week4# echo $?
0
```

Figure 8: Linux command './file1' to view the program output and 'echo $?' to view successful execution of the program.

From figure 8 we can see our Linux system on 64-bit architecture allocates the same space of 1, 4, 4, 8, 8, 8 byte for char, int, float, double, long, pointer respectively which matches the 64-bit column of Table I.

Question: Why might sizeof (long) differ between 32-bit and 64-bit systems?
Ans: "sizeof(long)" differs because 64-bit systems use larger word sizes to support bigger memory addresses, so "long" expands from 4 bytes to 8 bytes.

ii) Lab 2: Memory Segment Address Mapping

```
root@DESKTOP-OAP8OK0:/batch38/Week4# touch file2.c
root@DESKTOP-OAP8OK0:/batch38/Week4# nano file2.c
```

Figure 9: Linux command 'touch file2.c' to create file2.c and 'nano file2.c' to write C program inside file2.c.

```
root@DESKTOP-OAP8OK0:/batch38/Week4# cat file2.c
#include <stdio.h>
#include <stdlib.h>

int global_var = 150; // Data segment
int bss_var; // BSS/Data segment

int main() {
        int local_var = 30; // Stack
        static int local_static = 20; // Data segment

        int *heap_var = (int *)malloc(sizeof(int)); //Heap*heap_var = 500;
        printf(".......Variable Addresses.......\n\n");
        printf("global_var: %p (Data/BSS)\n", (void*)&global_var);
        printf("bss_var: %p (Data/BSS)\n", (void*)&bss_var);
        printf("local_var: %p (Stack)\n", (void*)&local_var);
        printf("local_static: %p (Data/BSS)\n", (void*)&local_static);
        printf("heap_var ptr: %p (Stack - pointer)\n",(void*)&heap_var);
        printf("*heap_var: %p (Heap - data)\n", (void*)heap_var);
        long diff = (long)((char*)heap_var - (char*)&local_var);
        printf("Address difference between stack and heap: %ld bytes\n", diff);

        free(heap_var);
        return 0;
}
```

Figure 10: Linux command 'cat file2.c' to view the contents.

```
root@DESKTOP-OAP8OK0:/batch38/Week4# gcc -E file2.c -o file2.i
root@DESKTOP-OAP8OK0:/batch38/Week4# gcc -S file2.i -o file2.s
root@DESKTOP-OAP8OK0:/batch38/Week4# gcc -c file2.s -o file2.o
root@DESKTOP-OAP8OK0:/batch38/Week4# gcc file2.o -o file2
```

Figure 11: Linux command of 4 compilation stages to convert file2.c to executable file "file2".

```
root@DESKTOP-OAP8OK0:/batch38/Week4# ./file2
.......Variable Addresses.......

global_var: 0x5cff4d08f010 (Data/BSS)
bss_var: 0x5cff4d08f01c (Data/BSS)
local_var: 0x7ffe540181b4 (Stack)
local_static: 0x5cff4d08f014 (Data/BSS)
heap_var ptr: 0x7ffe540181b8 (Stack - pointer)
*heap_var: 0x5cff686e6010 (Heap - data)
Address difference between stack and heap: -38478269325732 bytes
root@DESKTOP-OAP8OK0:/batch38/Week4# echo $?
0
```

Figure 12: Linux command './file2' to view the program output and 'echo $?' to view successful execution of the program.

```
root@DESKTOP-OAP8OK0:/batch38/Week4# touch file3.c
root@DESKTOP-OAP8OK0:/batch38/Week4# nano file3.c
```

Figure 13: Linux command 'touch file3.c' to create file3.c and 'nano file3.c' to write contents inside file3.c.

```
root@DESKTOP-OAP8OK0:/batch38/Week4# cat file3.c
#include <stdio.h>
#include <stdlib.h>

int main() {
        int *forheap; // Declare pointer (lives on STACK)
        forheap = (int *)malloc(sizeof(int)); //Allocate memory on HEAP first!
        *forheap = 30; // store value to allocated memory// Print addresses
        printf("Address of pointer (STACK): %p\n", (void*)&forheap);
        printf("Address of data (HEAP): %p\n", (void*)forheap);
        printf("Value stored: %d\n", *forheap);

        return 0;
}
```

Figure 14: Linux command 'cat file3.c' to view the contents inside file3.c.

```
root@DESKTOP-OAP8OK0:/batch38/Week4# gcc -E file3.c -o file3.i
root@DESKTOP-OAP8OK0:/batch38/Week4# gcc -S file3.i -o file3.s
root@DESKTOP-OAP8OK0:/batch38/Week4# gcc -c file3.s -o file3.o
root@DESKTOP-OAP8OK0:/batch38/Week4# gcc file3.o -o file3
```

Figure 15: Linux command of 4 compilation stages to convert "file3.c" to executable file "file3".

```
root@DESKTOP-OAP8OK0:/batch38/Week4# ./file3
Address of pointer (STACK): 0x7ffc469e7f30
Address of data (HEAP): 0x5fdd24cc1010
Value stored: 30
root@DESKTOP-OAP8OK0:/batch38/Week4# echo $?
0
root@DESKTOP-OAP8OK0:/batch38/Week4# ls
file1    file1.i  file1.s  file2.c  file2.o  file3    file3.i  file3.s
file1.c  file1.o  file2    file2.i  file2.s  file3.c  file3.o
```

Figure 16: Linux command './file3' to view the program output and 'echo $?' to view confirms successful execution of the program and 'ls' to view a list of files.

The output I observed from "file2 executable file" is that the program prints both the segment type and the hex address for each variable.I also observed that the address difference between stack and heap gives us a negative value of -38478269325732 bytes which proves stack address is numerically higher than the

heap. We can also verify the stack is above heap manually by comparing the hex value output of “figure 12” where we can see local_var (Stack) = 0x7ffe540181b4, and *heap_var (Heap) = 0x5cff686e6010, this tells us that 0x7ffe540181b4 > 0x5cff686e6010 because of hex value, and larger hex means higher memory address, therefore the stack is at top. And in “file3 executable file” I observed the pointer variable itself is on the stack, but the memory it points to (where “30” is stored) is in the heap.

Summary and Key Takeaways

Data types are contracts — they tell the OS how much memory to allocate and how to interpret it.

• Memory is segmented — stack for locals, heap for dynamic, data/BSS for globals, text for code. (we will further explore the text section in more details if necessary)

• The OS manages it all — tracking, allocating, protecting, and virtualizing memory.

• Stack is automatic; heap is manual — choose based on lifetime and size needs.

• Virtual memory provides isolation — each process has its own address space.

• Pointers bridge stack and heap — the pointer variable lives on stack, but points to heap memory.

Knowledge Check Question with Answers

1. What is a data type in the context of programming and operating systems?

Ans: A data type is a contract that tells the OS how many bytes to allocate, how to interpret those bytes, and what operations are valid.

---

2. Name the five main memory segments in a process memory layout.

Ans: Text, Data, BSS, Heap, Stack. (Here the order is placed from lower to higher address)

---

3. What is the difference between the Data segment and the BSS segment?

Ans: Data segment stores initialized global/static variables(value assigned), while BSS stores uninitialized global/static variables(values not assigned).

---

4. Why does the OS place the stack at high virtual addresses and the text segment at low addresses?

Ans: The stack is placed at high virtual addresses because in the process memory layout (also known as virtual memory map arrangement process) the OS reserves the highest address range(highest hex value) for the stack so that as new stack frames (i.e local variables and function calls) are added, the addresses decrease — so the stack "grows downward" numerically. (meaning it's hex value decreases).

The heap starts lower and grows upward when you allocate memory dynamically with "malloc" or similar functions.(here the heap's hex value increases for every new heap frames added)

By placing them at opposite ends, the OS ensures they can expand toward each other without colliding.

The text/code segment is fixed at low addresses so instructions remain stable, and the OS enforces boundaries to protect them from accidental overwrites. If we try to write/ overwrite on the text/code segment this will crash with a segmentation fault, because the OS has marked the text segment as read-only. That's the protection in action.

---

5. What happens to stack memory when a function returns? Who reclaims it?

Ans: The stack frame is automatically popped, and the compiler + OS reclaim the memory.

---

6. How does the OS know which memory belongs to which process?

Ans: Through virtual memory and page tables each process has its own address space mapped by the OS memory manager.

Github link: OS-Security-And-Networks/Week4/Lab4 Task.pdf at main · firstnamecheten/OS-Security-And-Networks